

## Problem Statement or Requirement:

A requirement from the Hospital, Management asked us to create a predictive model which will predict the Chronic Kidney Disease (CKD) based on the several parameters. The Client has provided the dataset of the same

### CKD Prediction – Model Development Report

#### 1) Problem statement

Build a supervised ML model that predicts **Chronic Kidney Disease (CKD)** (yes/no) from routine clinical and lab parameters to enable early triage and follow-up.

#### 2) Dataset overview

- **Rows × columns:** 399 × 25
- **Target:** classification (yes/no)
- **Class distribution:** yes=249 (62.4%), no=150 (37.6%) → mild imbalance

```
dataset=pd.read_csv("CKD.csv")
dataset['classification'].value_counts()

classification
yes    249
no     150
Name: count, dtype: int64
```

[4]:	age	bp	al	su	bgr	bu	sc	sod	pot	hrmo	...	pc_normal	pcc_present	ba_present	ht
0	2.000000	76.459948	3.0	0.0	148.112676	57.482105	3.077356	137.528754	4.627244	12.518156	...	0	0	0	
1	3.000000	76.459948	2.0	0.0	148.112676	22.000000	0.700000	137.528754	4.627244	10.700000	...	1	0	0	
2	4.000000	76.459948	1.0	0.0	99.000000	23.000000	0.600000	138.000000	4.400000	12.000000	...	1	0	0	
3	5.000000	76.459948	1.0	0.0	148.112676	16.000000	0.700000	138.000000	3.200000	8.100000	...	1	0	0	
4	5.000000	50.000000	0.0	0.0	148.112676	25.000000	0.600000	137.528754	4.627244	11.800000	...	1	0	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
394	51.492308	70.000000	0.0	0.0	219.000000	36.000000	1.300000	139.000000	3.700000	12.500000	...	1	0	0	
395	51.492308	70.000000	0.0	2.0	220.000000	68.000000	2.800000	137.528754	4.627244	8.700000	...	1	0	0	
396	51.492308	70.000000	3.0	0.0	110.000000	115.000000	6.000000	134.000000	2.700000	9.100000	...	1	0	0	
397	51.492308	90.000000	0.0	0.0	207.000000	80.000000	6.800000	142.000000	5.500000	8.500000	...	1	0	0	
398	51.492308	80.000000	0.0	0.0	100.000000	49.000000	1.000000	140.000000	5.000000	16.300000	...	1	0	0	

399 rows × 28 columns

#### 3) Data preprocessing

- **Missing tokens → NaN:** replaced ? / NA / na / NaN / nan / "" with NaN.

- **Imputation:**
  - **Numeric:** median (fit on train, applied to test)
  - **Categorical:** most-frequent value (fit on train, applied to test)
- **Encoding:** one-hot encoding (train on train set; **test columns aligned** to train to handle unseen categories).
- **Scaling:**
  - Applied **StandardScaler** for models that benefit from scaling (**Logistic Regression, kNN, GaussianNB, SVM**).
  - **Not** applied for tree/boosting models (**Random Forest, Decision Tree**).
- **Evaluation protocol:** **Stratified** train/test split (test size  $\approx 1/3 \rightarrow 133$  records). Model selection via **5-fold stratified CV**.  
Primary metric: **ROC AUC**; also report Accuracy, Precision, Recall, F1, PR-AUC on the hold-out test set.

#### 4) Experiments & results

##### Models evaluated

- SVM (RBF/Linear grid)
- Logistic Regression
- Random Forest
- Decision Tree
- k-Nearest Neighbors (kNN)
- Gaussian Naive Bayes
- Xgboost

	A	B	C	D	E	F	G	H	I
1	<b>CKD Prediction – Model Development Report</b>								
2									
3	<b>Model</b>	<b>Best CV R</b>	<b>Best Params (key)</b>	<b>Test Accu</b>	<b>Precision</b>	<b>Recall (pos)</b>	<b>F1 (pos)</b>	<b>ROC-AUC</b>	<b>Test PR-AUC</b>
4	SVM (grid)	—	kernel={rbf,linear}, C∈{0.5,1,3,10}, γ∈{scale,0.1,0.01}	0.98	1	0.98	0.99	0.988	—
5	Logistic Regression	1	C=0.1, solver=lbfgs, penalty=l2	0.970	1.000	0.952	0.975	1.000	1.000
6	Random Forest	1	n_estimators=200, max_depth=None, max_features=√, min_samples_split=2, min_samples_leaf=1	0.985	0.988	0.988	0.988	1.000	1.000
7	Decision Tree	0.971	criterion=gini, max_depth=None, min_samples_split=10, min_samples_leaf=2	0.970	0.976	0.976	0.976	0.976	0.977
8	kNN	1	n_neighbors=11, p=1, weights=uniform	0.932	1.000	0.892	0.943	0.993	0.995
9	GaussianNB	1	var_smoothing=1e-12	0.977	0.988	0.976	0.982	0.990	0.988
10	Xgboost		colsample_bytree: 0.8, learning_rate: 0.05,	0.925	0.958	0.920	0.939	0.992	0.995
11									
12									
13									

SVM(RBF)

Fitting 5 folds for each of 24 candidates, totalling 120 fits				
[[49 1]				
[ 1 82]]				
	precision	recall	f1-score	support
0	0.98	0.98	0.98	50
1	0.99	0.99	0.99	83
accuracy			0.98	133
macro avg	0.98	0.98	0.98	133
weighted avg	0.98	0.98	0.98	133
Accuracy : 0.9849624060150376				
Precision: 0.9879518072289156				
Recall : 0.9879518072289156				
F1 : 0.9879518072289156				
ROC AUC : 0.9990361445783132				
Best params: {'C': 0.5, 'gamma': 'scale', 'kernel': 'rbf'}				

- Logistic Regression

Fitting 5 folds for each of 26 candidates, totalling 130 fits

==== LogisticRegression =====

Best CV score ( roc\_auc ): 1.0

Best params: {'C': np.float64(0.1), 'penalty': 'l2', 'solver': 'lbfgs'}

Confusion matrix:

```
[[50  0]
 [ 4 79]]
```

Classification report:

	precision	recall	f1-score	support
0	0.93	1.00	0.96	50
1	1.00	0.95	0.98	83
accuracy			0.97	133
macro avg	0.96	0.98	0.97	133
weighted avg	0.97	0.97	0.97	133

Accuracy : 0.9699248120300752

Precision: 1.0

Recall : 0.9518072289156626

F1 : 0.9753086419753086

ROC AUC : 0.9995180722891566

PR AUC : 0.9997165131112689

Fitting 5 folds for each of 324 candidates, totalling 1620 fits

- Random Forest

Fitting 5 folds for each of 324 candidates, totalling 1620 fits

==== RandomForest =====

Best CV score ( roc\_auc ): 1.0

Best params: {'max\_depth': None, 'max\_features': 'sqrt', 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 200}

Confusion matrix:

```
[[49  1]
 [ 1 82]]
```

Classification report:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	50
1	0.99	0.99	0.99	83
accuracy			0.98	133
macro avg	0.98	0.98	0.98	133
weighted avg	0.98	0.98	0.98	133

Accuracy : 0.9849624060150376

Precision: 0.9879518072289156

Recall : 0.9879518072289156

F1 : 0.9879518072289156

ROC AUC : 0.9995180722891566

PR AUC : 0.9997114101846285

- Decision Tree

Fitting 5 folds for each of 162 candidates, totalling 810 fits

==== DecisionTree =====

Best CV score ( roc\_auc ): 0.9711051693404634

Best params: {'criterion': 'gini', 'max\_depth': None, 'min\_samples\_leaf': 2, 'min\_samples\_split': 10}

Confusion matrix:

```
[[48  2]
 [ 2 81]]
```

Classification report:

	precision	recall	f1-score	support
0	0.96	0.96	0.96	50
1	0.98	0.98	0.98	83
accuracy			0.97	133
macro avg	0.97	0.97	0.97	133
weighted avg	0.97	0.97	0.97	133

Accuracy : 0.9699248120300752

Precision: 0.9759036144578314

Recall : 0.9759036144578314

F1 : 0.9759036144578314

ROC AUC : 0.9756626506024098

PR AUC : 0.9771266510446963

- k-Nearest Neighbors (kNN)

Fitting 5 folds for each of 60 candidates, totalling 300 fits

===== kNN =====

Best CV score ( roc\_auc ): 1.0

Best params: {'n\_neighbors': 11, 'p': 1, 'weights': 'uniform'}

Confusion matrix:

```
[[50  0]
 [ 9 74]]
```

Classification report:

	precision	recall	f1-score	support
0	0.85	1.00	0.92	50
1	1.00	0.89	0.94	83
accuracy			0.93	133
macro avg	0.92	0.95	0.93	133
weighted avg	0.94	0.93	0.93	133

Accuracy : 0.9323308270676691

Precision: 1.0

Recall : 0.891566265060241

F1 : 0.9426751592356688

ROC AUC : 0.993012048192771

PR AUC : 0.9948934245906536

- Gaussian Naive Bayes

Fitting 5 folds for each of 7 candidates, totalling 35 fits

==== GaussianNB =====

Best CV score ( roc\_auc ): 1.0

Best params: {'var\_smoothing': np.float64(1e-12)}

Confusion matrix:

```
[[49  1]
 [ 2 81]]
```

Classification report:

	precision	recall	f1-score	support
0	0.96	0.98	0.97	50
1	0.99	0.98	0.98	83
accuracy			0.98	133
macro avg	0.97	0.98	0.98	133
weighted avg	0.98	0.98	0.98	133

Accuracy : 0.9774436090225563

Precision: 0.9878048780487805

Recall : 0.9759036144578314

F1 : 0.9818181818181818

ROC AUC : 0.9897590361445783

PR AUC : 0.9878101465936793

- Xgboost

Fitting 5 folds for each of 32 candidates, totalling 160 fits

==== XGBoost (gbtree, small grid) =====

Best CV score ( roc\_auc ): 0.9987394957983193

Best params: {'colsample\_bytree': 0.8, 'learning\_rate': 0.05, 'max\_depth': 3, 'min\_child\_weight': 1, 'subsample': 0.8}

Confusion matrix:

```
[[14  1]
 [ 2 23]]
```

Classification report:

	precision	recall	f1-score	support
0	0.88	0.93	0.90	15
1	0.96	0.92	0.94	25
accuracy			0.93	40
macro avg	0.92	0.93	0.92	40
weighted avg	0.93	0.93	0.93	40

Accuracy : 0.925

Precision: 0.9583333333333334

Recall : 0.92

F1 : 0.9387755102040817

ROC AUC : 0.992

PR AUC : 0.995194871794872

## 6) Final model & justification

Selected model: Random Forest Classifier

Best params (from grid): **n\_estimators=200**, **max\_depth=None**, **max\_features='sqrt'**,  
**min\_samples\_split=2**, **min\_samples\_leaf=1** (with **class\_weight='balanced'**)

Test performance (hold-out):

Accuracy: 0.985

Precision (pos): 0.988

Recall (pos): 0.988

F1 (pos): 0.988

ROC-AUC / PR-AUC: ~0.9995 / 0.9997

Confusion matrix: [[49, 1], [1, 82]] → only 1 false negative, which is critical for CKD screening.

Why this model:

Best overall balance of Precision, Recall, and F1 on the test set, with near-perfect AUCs.

Low clinical risk: extremely few false negatives (missed CKD), while keeping false positives low.

Stable & robust: 5-fold CV ROC-AUC  $\approx$  1.0, minimal variance across folds.

Why not the others:

SVM: strong (Acc 0.98, Rec 0.98) but more FN than RF (2 vs 1), and needs feature scaling + more tuning.

Logistic Regression: excellent AUC but lower recall (0.952) vs RF, risking more missed CKD cases.

Decision Tree: interpretable but weaker overall metrics; can overfit as a single tree.

kNN: lower recall (0.892) and accuracy on this dataset.

GaussianNB: very good, but slightly behind RF on recall/F1.

XGBoost (quick config tested): high AUC, but the tested setup underperformed RF on accuracy/recall; RF gave the best balance here.

- Balanced errors (confusion matrix [[49, 1], [1, 82]])—very low false negatives, which is crucial for CKD screening.